PRODUCT REQUIREMENTS DOCUMENT

WorkRadar

One AI-ranked view of a team's work across Slack, Jira, GitHub and email, for Airtribe's engineering managers.

Aditya Singh
Airtribe · AI Product Management
v1.1 · 21 October 2026
Live: **work-radar-ten.vercel.app**

01 Problem statement

An engineering manager on Airtribe's Learning Platform team starts every day by rebuilding the state of the team from four tools. The signals that matter are a P0 bug stopping learners from joining tonight's live class, a designer blocked for two days, and a launch go/no-go waiting on the manager's reply. They sit beside hundreds of FYIs and bot updates, and each tool marks priority differently, or not at all. Triage eats the first hour, blockers surface a day late, and escalations get missed.

> **How might we give a manager one trustworthy, prioritised answer to "what needs me today?", so they can act in five minutes instead of searching for an hour?**

WorkRadar ranks every item with a transparent score, has Gemini turn the ranking into a daily brief, and lets the manager assign, finish or snooze work without switching tools.

A manager's day today (estimates to validate)

4 tools to check before acting	**~250** updates a day: Slack, Jira, PRs, mail	**60–90m** daily triage and status-chasing
<5% of updates need the manager today	**~1 day** before a blocker reaches the manager	**0** shared view of "waiting on me"

02 Target users

Primary: engineering managers and tech leads at Airtribe with 5–12 reports, such as Learning Platform, Growth and Data. Their work runs to fixed external deadlines: live classes, submission deadlines and launches.

Secondary: program and product leaders who coordinate cross-team launches like the GenAI Launchpad, and who need status and decisions without pinging the manager.

Not for v1: individual contributors, and directors who need cross-team roll-ups.

03 User personas

NK — Neha Kapoor, 34 · primary
Engineering Manager, Learning Platform · 8 reports

Goals Protect live classes, unblock her team fast, ship the GenAI Launchpad on time.

Frustrations Constant context switching. Requests to her get buried in threads, and she hears about blockers at stand-up.

Today Slack, Jira, GitHub, Gmail and a paper to-do list. About 75 minutes of triage before her first decision.

"By the time I've read everything, the day has picked my priorities for me."

RB — Rohan Bhatia, 41 · secondary
Head of Product · launch stakeholder

Goals Make fast, confident go/no-go calls, and get launch status he can trust without a meeting.

Frustrations Chasing managers for answers. Surprises on launch day.

Needs His asks visible and correctly ranked on each manager's radar, with clear owners.

"I don't need every ticket. I need to know nothing is stuck on us."

04 User journey: Neha's Tuesday with WorkRadar

08:55 · Browser — Opens one tab
WorkRadar is her start page. In about 2 seconds Gemini's brief says 4 items need her, and to start with the iPhone live-class bug.
Before: opens 4 tools and 60+ tabs

09:00 · Today — Sees the whole radar
Counts, a radar of 38 open items, the top 3 actions and what happened overnight, all on one screen.
Before: rebuilds the picture in her head

09:05 · Item detail — Acts with context
Reads why LP-512 scores 86, then assigns it to Arjun as the AI suggested. She snoozes an FYI. Every action can be undone.
Before: hunts for the owner and the thread

09:10 · Team — Unblocks people
Meera has been blocked on marketing copy for 2 days. The suggestion is to approve placeholder copy and nudge Tanvi's manager.
Before: learns this at tomorrow's stand-up

All day · Feed, ⌘K — Stays on top
Works the keyboard triage feed and searches "blocked prs waiting on me". Outages and AI failures degrade gracefully.
Before: notification whack-a-mole

05 Key pain points

Work is fragmented
Four tools, four priority systems. Jira has P-levels. Slack and email have none, yet hold the most urgent asks.

"Waiting on me" is invisible
Decisions, reviews and replies addressed to the manager get buried in threads and inboxes.

Blockers hide across tools
A PR waits on a ticket, which waits on another team. Nobody sees the chain until someone complains.

Urgency depends on time
A 7 pm live class or a same-day go/no-go matters more than a "P1" due next week. Tools don't rank by deadline.

Noise drowns signal
Bots, Dependabot, FYIs and CC chains make up most of the volume and train managers to skim.

Black-box AI isn't trusted
A summary with no reasons can't be acted on. Managers need to see why something is urgent.

06 Success metrics

targets 8 weeks after a 10-manager pilot

METRIC	TARGET	WHY IT MATTERS	HOW WE MEASURE IT
NORTH STAR **Time to clarity**	**< 5 min**	From an hour of triage to minutes	Session start → first action on a top-3 ranked item
Weekly active managers / invited	**≥ 70%**	Habit and lasting value	Managers with ≥ 3 active days in the week
Morning check-in	**≥ 60% of days**	The daily habit is the product	Today opened before 10:00 on a workday
Blocker time to resolution	**–30%**	Highest-leverage manager action	Blocked flag → cleared, against a 4-week baseline
Actions taken in WorkRadar	**≥ 5 / day**	People act here, not just read	Assign, done and snooze per manager, net of undos
Ranking trust	**≥ 80%**	An untrusted radar gets ignored	Weekly "top 3 were right" pulse; re-rank rate < 15%
GUARDRAIL Missed P0/P1 escalations	**0**	Costs learners and revenue	Every post-mortem checks it reached the inner ring
GUARDRAIL AI factual errors	**< 2% of briefs**	One invented fact breaks trust	Weekly eval of 50 briefs against their snapshots
GUARDRAIL Brief latency · cost	**p95 < 6 s · < ₹1**	Daily use at pilot scale	Vercel function logs and Gemini token usage

07 MVP scope

In scope (built and working in the demo)

- Four sources: Slack, Jira, GitHub, Gmail (realistic mock data shaped like each API)
- Today dashboard: radar, AI brief, needs attention, who's blocked
- Priority feed with an explainable 0–100 score and three rings
- Item detail, plain-language search, ⌘K palette, Team view
- Assign (AI-suggested owner), Mark done, Snooze, Undo
- Live Gemini brief through a serverless route, with on-device fallback
- Empty, error and outage states; dark mode; mobile layout
- Settings: AI engine, live ranking weights, outage simulation

Out of scope (v1)

- Live OAuth connectors and write-back to source tools
- Director roll-ups; weights learned from behaviour
- Push digests and alerts (Slack, email)
- Teams, Outlook, Linear, Notion; admin and SSO

08 Assumptions

1. **Explainable beats clever.** Managers will trust a deterministic score they can inspect and re-weight more than an opaque AI ranking. AI writes the brief; it doesn't rank.
2. **Read-only access is acceptable.** Airtribe security will approve least-privilege OAuth scopes if message bodies are kept no longer than 7 days and never used for training.
3. **A fast start page builds the habit.** If the brief is ready in seconds, managers will open WorkRadar first each morning. A Slack or email digest is the next lever if they don't.
4. **Time is the strongest signal.** Deadline proximity, who is waiting and what is blocked predict "needs me today" better than labels.
5. **The volume estimates hold.** To validate with 5 manager interviews and a 2-week diary study before the pilot.
6. **Cost is negligible.** Gemini Flash-Lite on Vercel's free tier keeps a brief under ₹1 and pilot running costs near zero.

09 Features

MUST-HAVE MVP, shipped in the demo

- **Unified priority feed** ranked by six weighted signals: urgency, impact, unblocks others, waiting on you, stuck, escalation.
- **AI daily brief**: top 3 for today, what happened, needs attention, who's blocked and how to unblock them.
- **"Why it's on your radar"**: score breakdown on every item, with adjustable weights.
- **Item detail**: AI summary, next step, dependencies and source context.
- **Natural-language search** ("blocked prs waiting on me") and ⌘K.
- **One-click actions**: Assign (AI-suggested owner), Done, Snooze, Undo.
- **Team blockers**, sorted by who needs the manager first.
- **Graceful degradation**: Flash-Lite → Flash → on-device, with the reason shown; outage banner.

NICE-TO-HAVE Next releases

- **Write-back**: assign in Jira, request reviews in GitHub, reply in Slack.
- **AI reply drafts** for asks waiting on the manager.
- **Learned weights** from what gets opened, snoozed or reordered.
- **Calendar awareness**: plan the top 3 around meetings and focus time.
- **1:1 prep and weekly digest** for each report: shipped, stuck, waiting.
- **Director roll-ups** across teams, with launch risk heatmaps.
- **More sources**: Linear, Notion, Teams, Outlook, PagerDuty.
- **AI triage of live threads**: priority, owner and next step for Slack and email.
- **Morning digest** in Slack or email, with quiet hours.

Rollout	Week 0: interviews and diary study → Weeks 1–2: pilot with 3 Learning Platform managers (read-only) → Weeks 3–8: 10 managers with write-back, weekly metric review → Go/no-go on org-wide rollout against the north star and guardrails.
Top risks	Over-trusting the AI (mitigated by explainable scores, labelled engines and undo) · Privacy of message content (read-only scopes, short retention) · Ranking drift as teams change (adjustable weights, weekly trust pulse).